UNIVERSITEIT VAN PRETORIA
UNIVERSITY OF PRETORIA
YUNIBESITHI YA PRETORIA

Faculty of Engineering, Built Environment and Information Technology

Fakulteit Ingenieurswese, Bou-omgewing en Inligtingtegnologie / Lefapha la Boetšenere, Tikologo ya Kago le Theknolotši ya Tshedimošo

Special Topic: Word Embeddings + Language Models

Dr. Vukosi Marivate

Inputs:
- A Modupe [PhD Candidate]
- A Moodley [MIT Big Data Science Student]

Make today matter

Data Science for Social Impact

Word Embeddings + Language Models

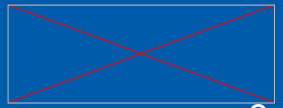

Word Embeddings

“You shall know a word by the company it keeps”

“Tell me who your friends are, and I will tell you who you are.”

The Distributional Hypothesis is that words that occur in the same contexts tend to have similar meanings [2]

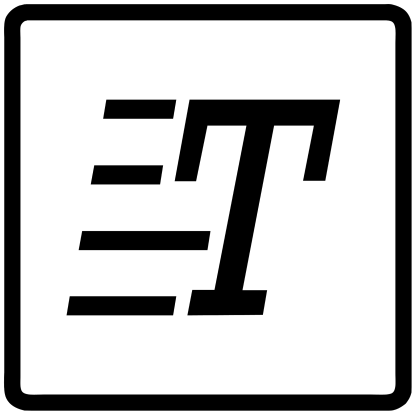

Word Embeddings

“You shall know a word by the company it keeps”

“Tell me who your friends are, and I will tell you who you are.”

The Distributional Hypothesis is that words that occur in the same contexts tend to have similar meanings [2]

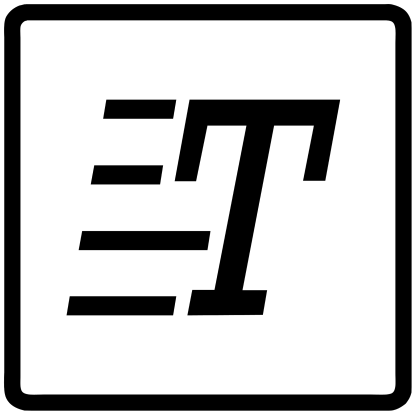

Word Vectors

- Mapping from tokens to a continuous vector space
- Trained using a shallow neural network (not deep)

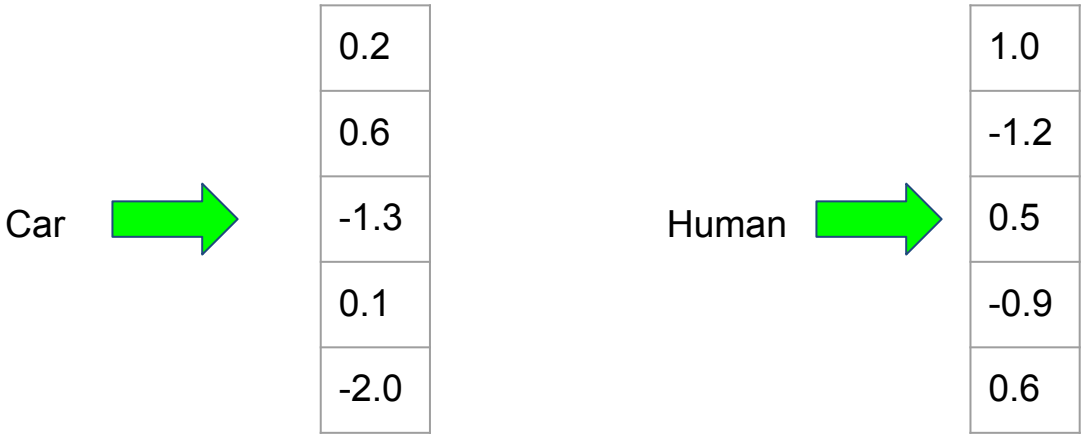

https://www.youtube.com/watch?v=UqRCEmrv1gQ

Word Vectors - The Idea!

- I went out with my **BOYFRIEND** to the movies.
- I went out with my **GIRLFRIEND** to the movies.
- I went out with my **BAE** to the movies.
- I went out with my **FRIENDS** to the movies.

Words used in a similar fashion in the same context!!!!

https://www.youtube.com/watch?v=UqRCEmrv1gQ

CBOW - Continuous Bag of Words

Predict word from context

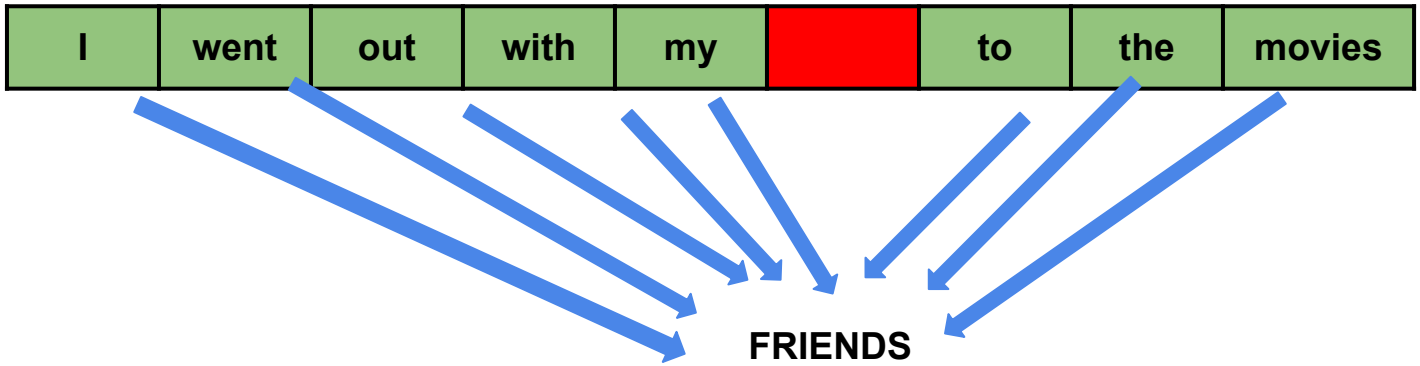

https://www.youtube.com/watch?v=UqRCEmrv1gQ

CBOW - Working

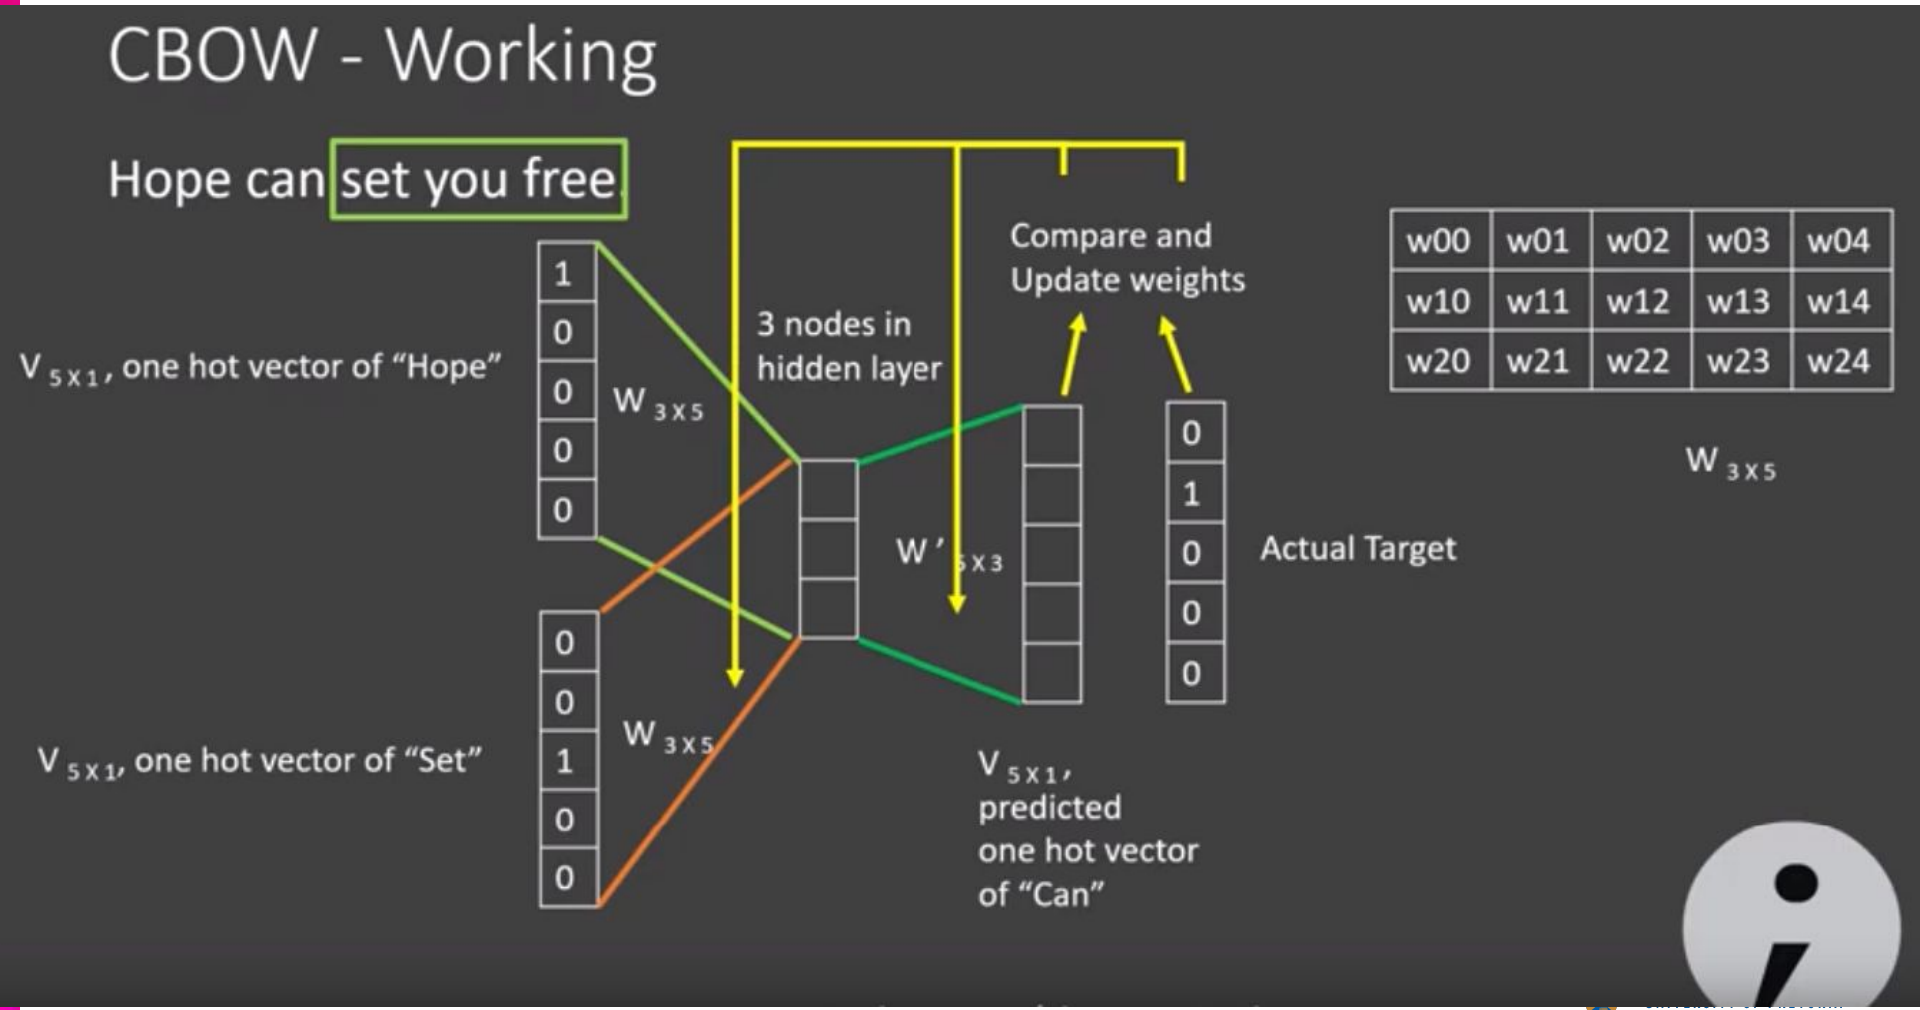

https://www.youtube.com/watch?v=UqRCEmrv1gQ

Skip-GRAM: Crazy idea, but works!!

Predict context from words

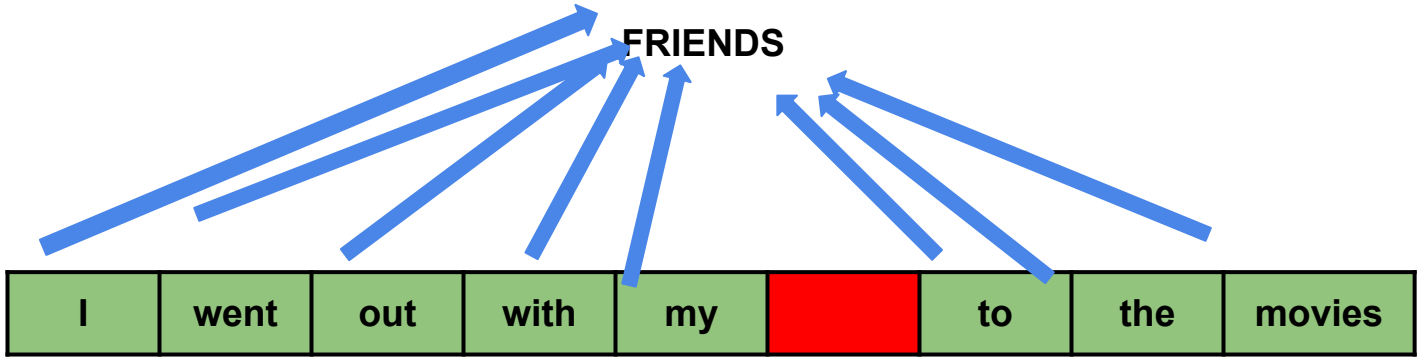

https://www.youtube.com/watch?v=UqRCEmrv1gQ

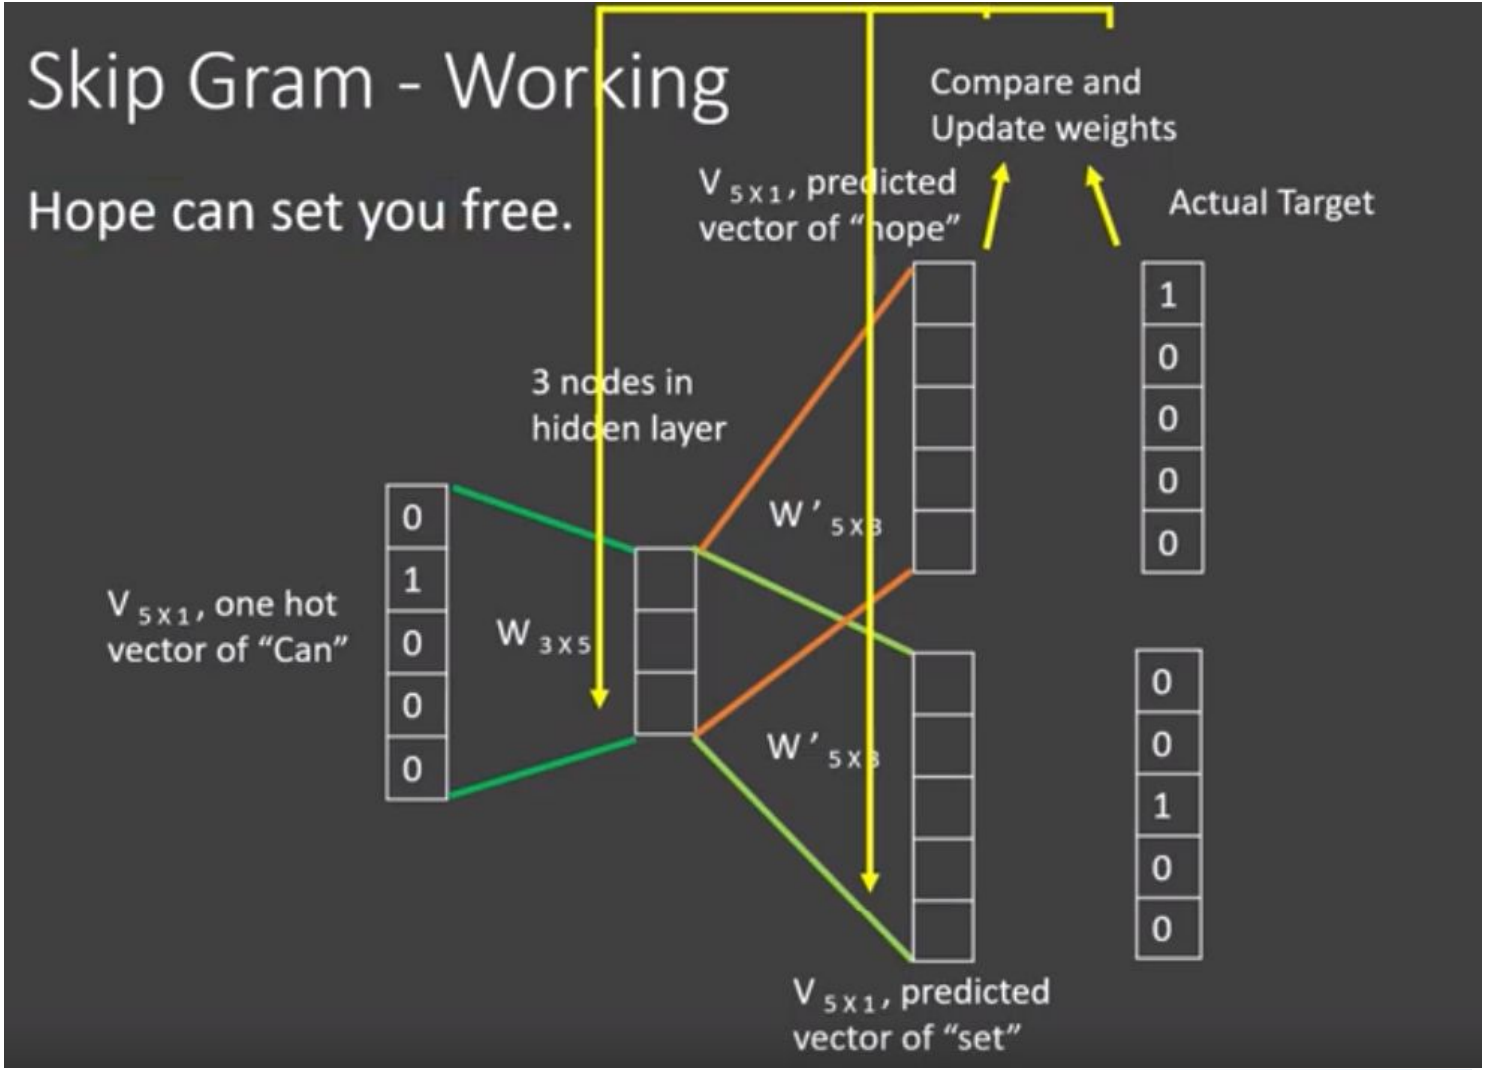

https://www.youtube.com/watch?v=UqRCEmrv1gQ

Word2Vec

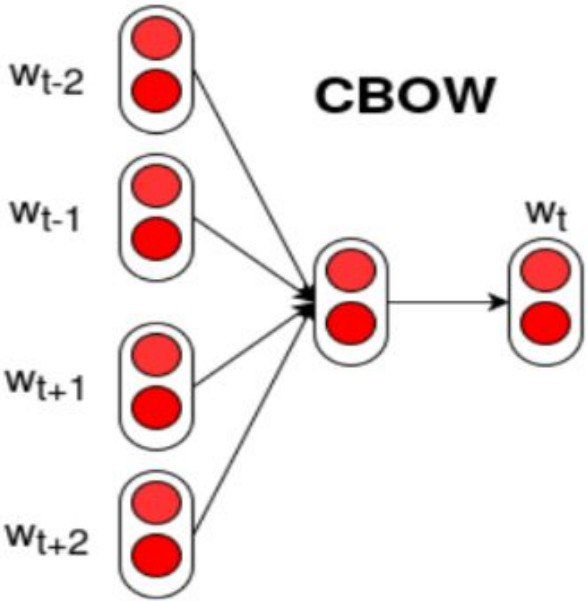

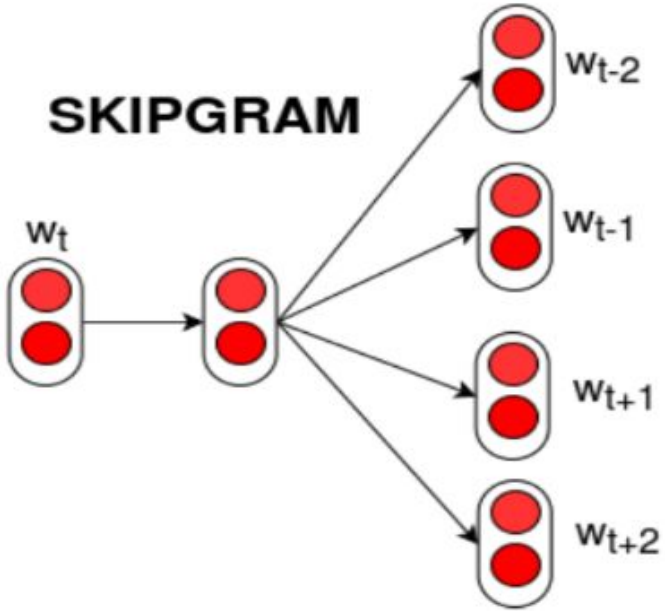

$$\frac{1}{T}\sum_{t=1}^{T}\sum_{-c\leq j\leq c, j\neq 0} logp(w_t|w_{t+j})$$

$$\frac{1}{T}\sum_{t=1}^{T}\sum_{-c\leq j\leq c, j\neq 0} logp(w_{t+j}|w_t)$$

https://arxiv.org/abs/1301.3781 [Available in Gensim to train], Pretrained as well

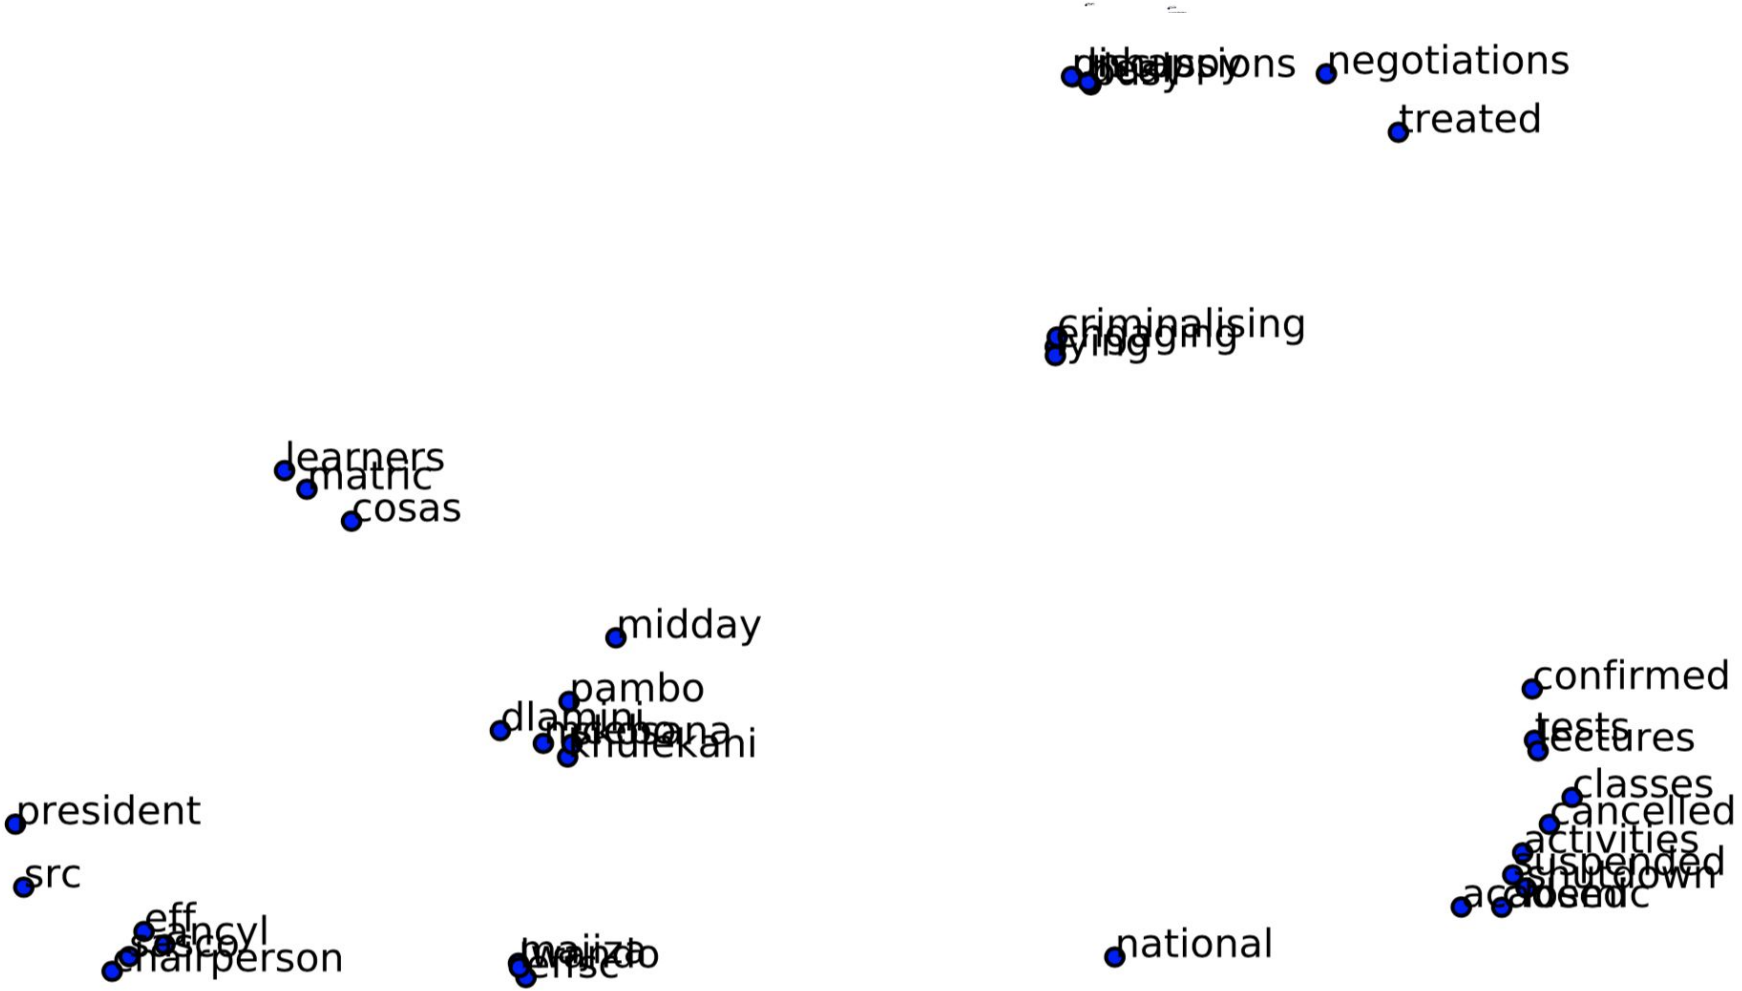

Fees Must Fall Twitter Data. Word2vec + TSNE

GloVe: Global Vectors for Word Representation

Nearest neighbors

The Euclidean distance (or cosine similarity) between two word vectors provides an effective method for measuring the linguistic or semantic similarity of the corresponding words. Sometimes, the nearest neighbors according to this metric reveal rare but relevant words that lie outside an average human's vocabulary. For example, here are the closest words to the target word *frog*:

0. *frog*
1. frogs
2. toad
3. litoria
4. leptodactylidae
5. rana
6. lizard
7. eleutherodactylus

3. litoria

4. leptodactylidae

5. rana

7. eleutherodactylus

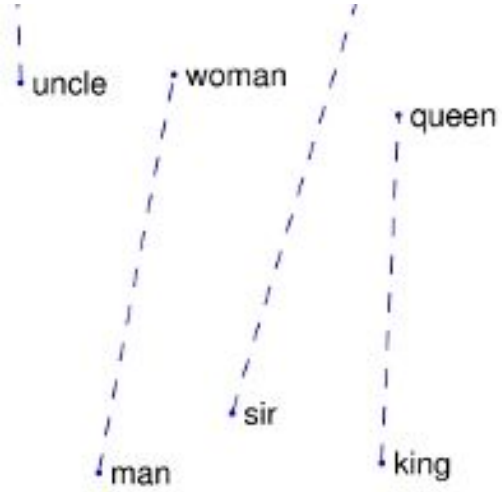

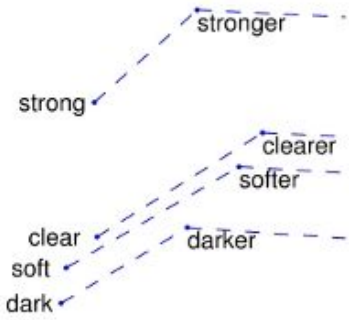

Linear Substructures: Analogies

https://nlp.stanford.edu/projects/glove/ [Pre built vectors]

Downstream tasks: Classification

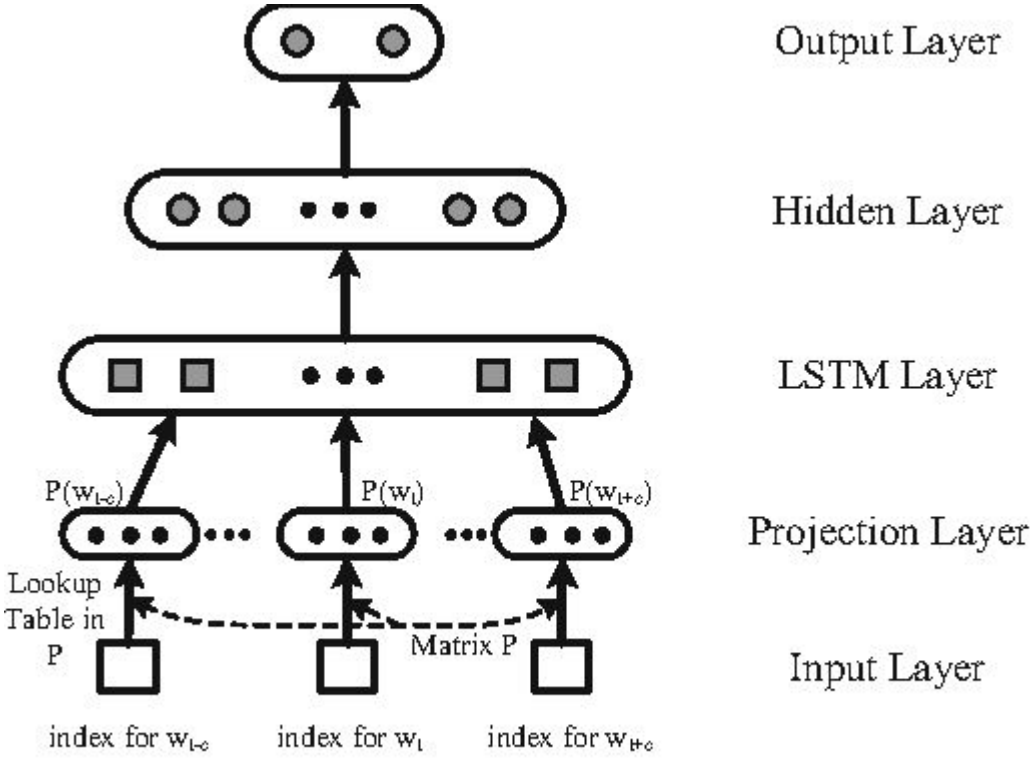

Advanced: Language Models

AI researchers debate the ethics of sharing potentially harmful programs

Nonprofit lab OpenAI withheld its latest research, but was criticized by others in the field

By James Vincent | Feb 21, 2019, 10:30am EST

SHARE

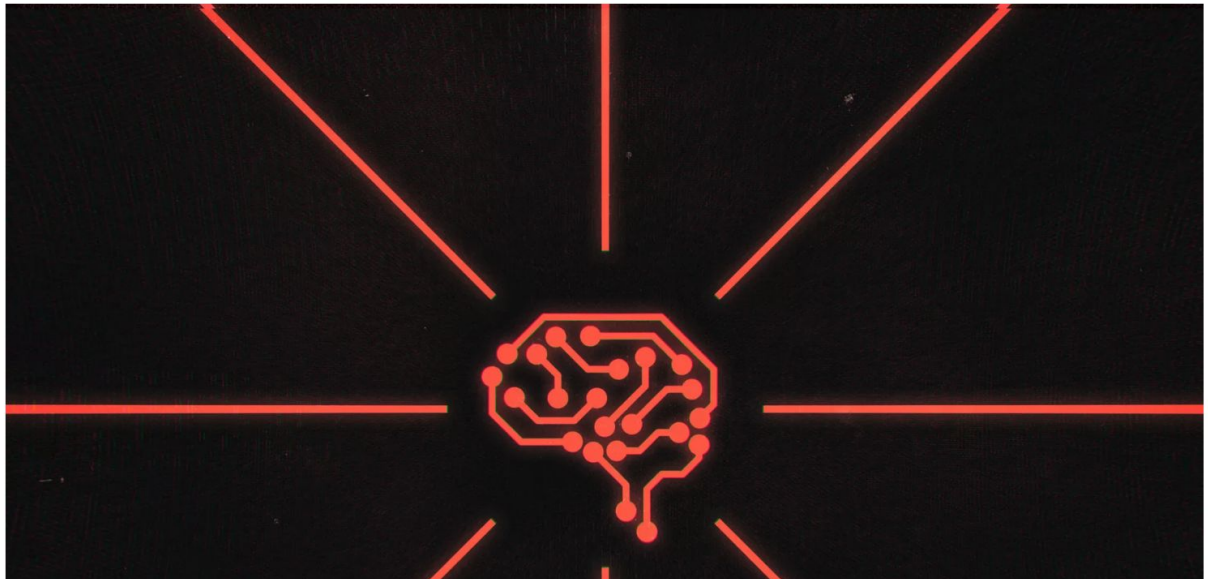

Beats Powerb
headphones a

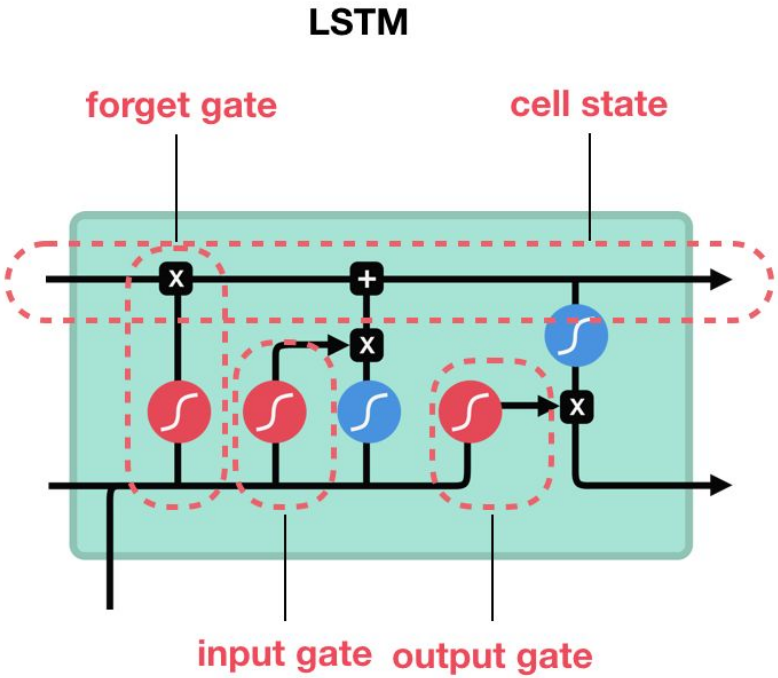

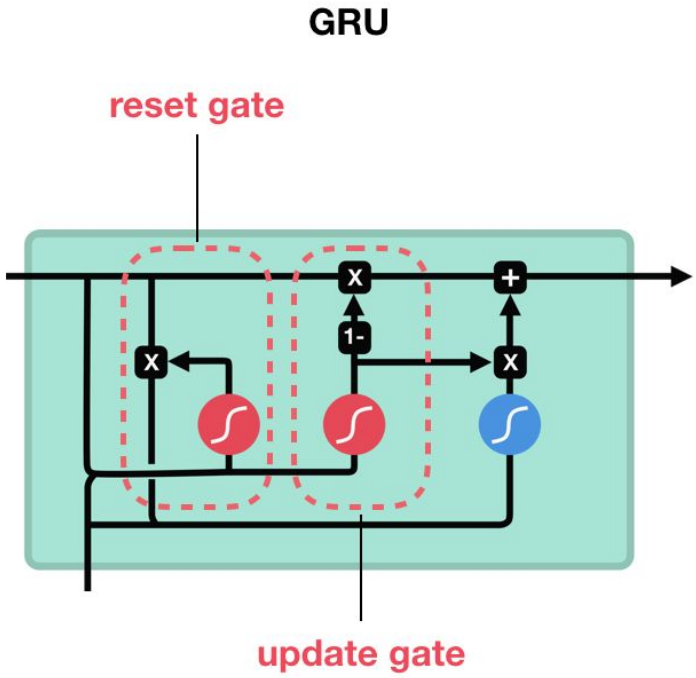


sigmoid


tanh


pointwise multiplication


pointwise addition

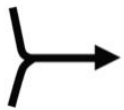
vector concatenation

Illustrated Guide to LSTM's and GRU's: A step by step explanation

Sequence Models

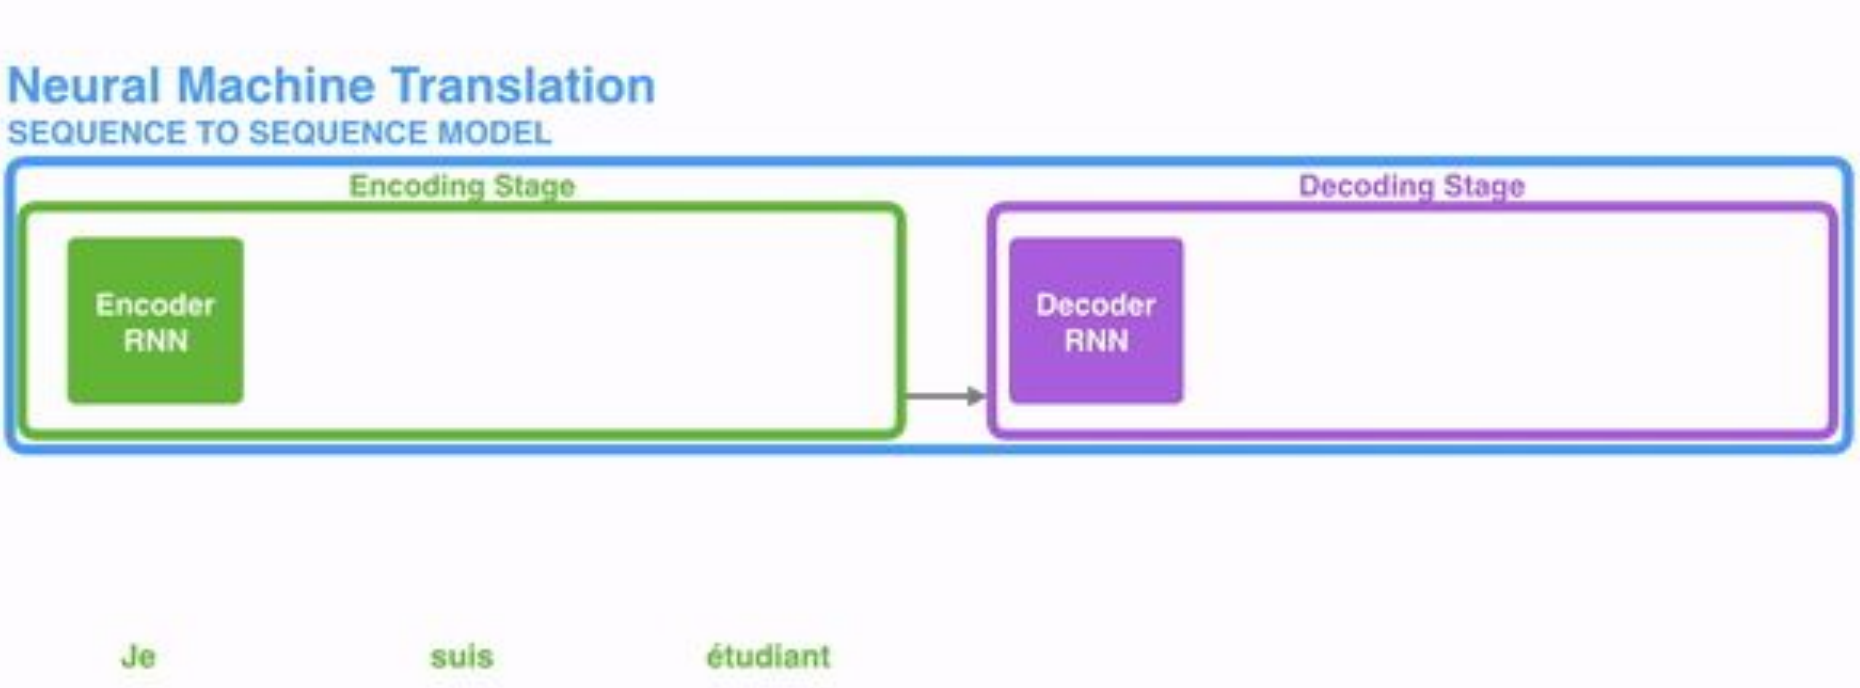

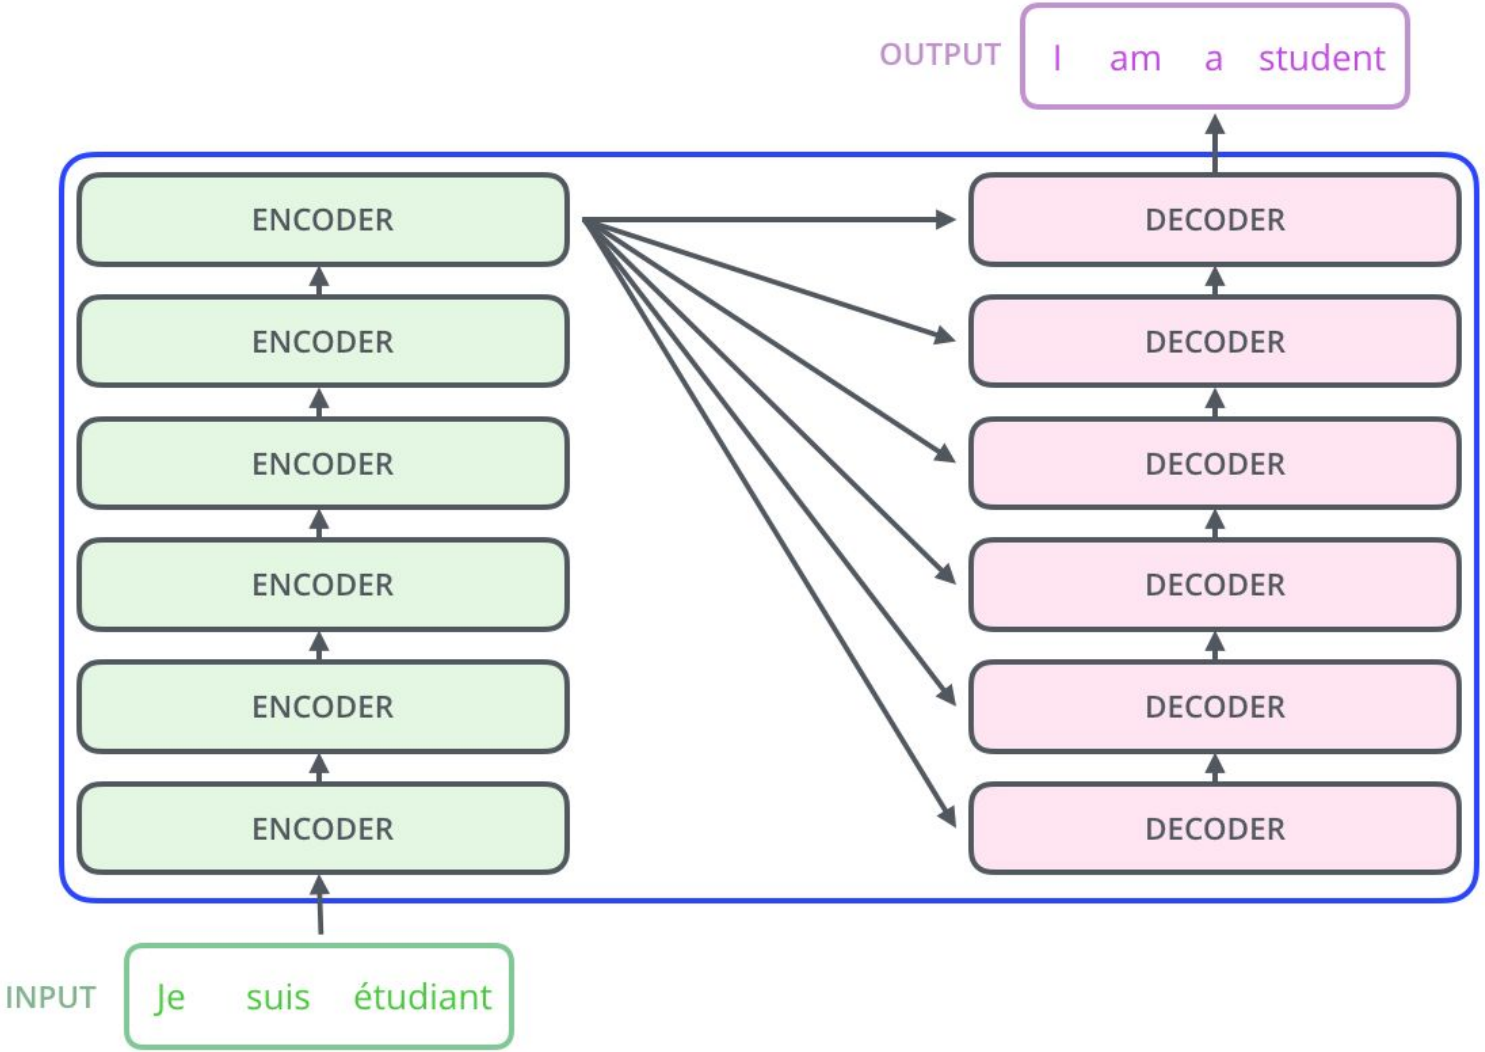

The Illustrated Transformer: http://jalammar.github.io/illustrated-transformer/

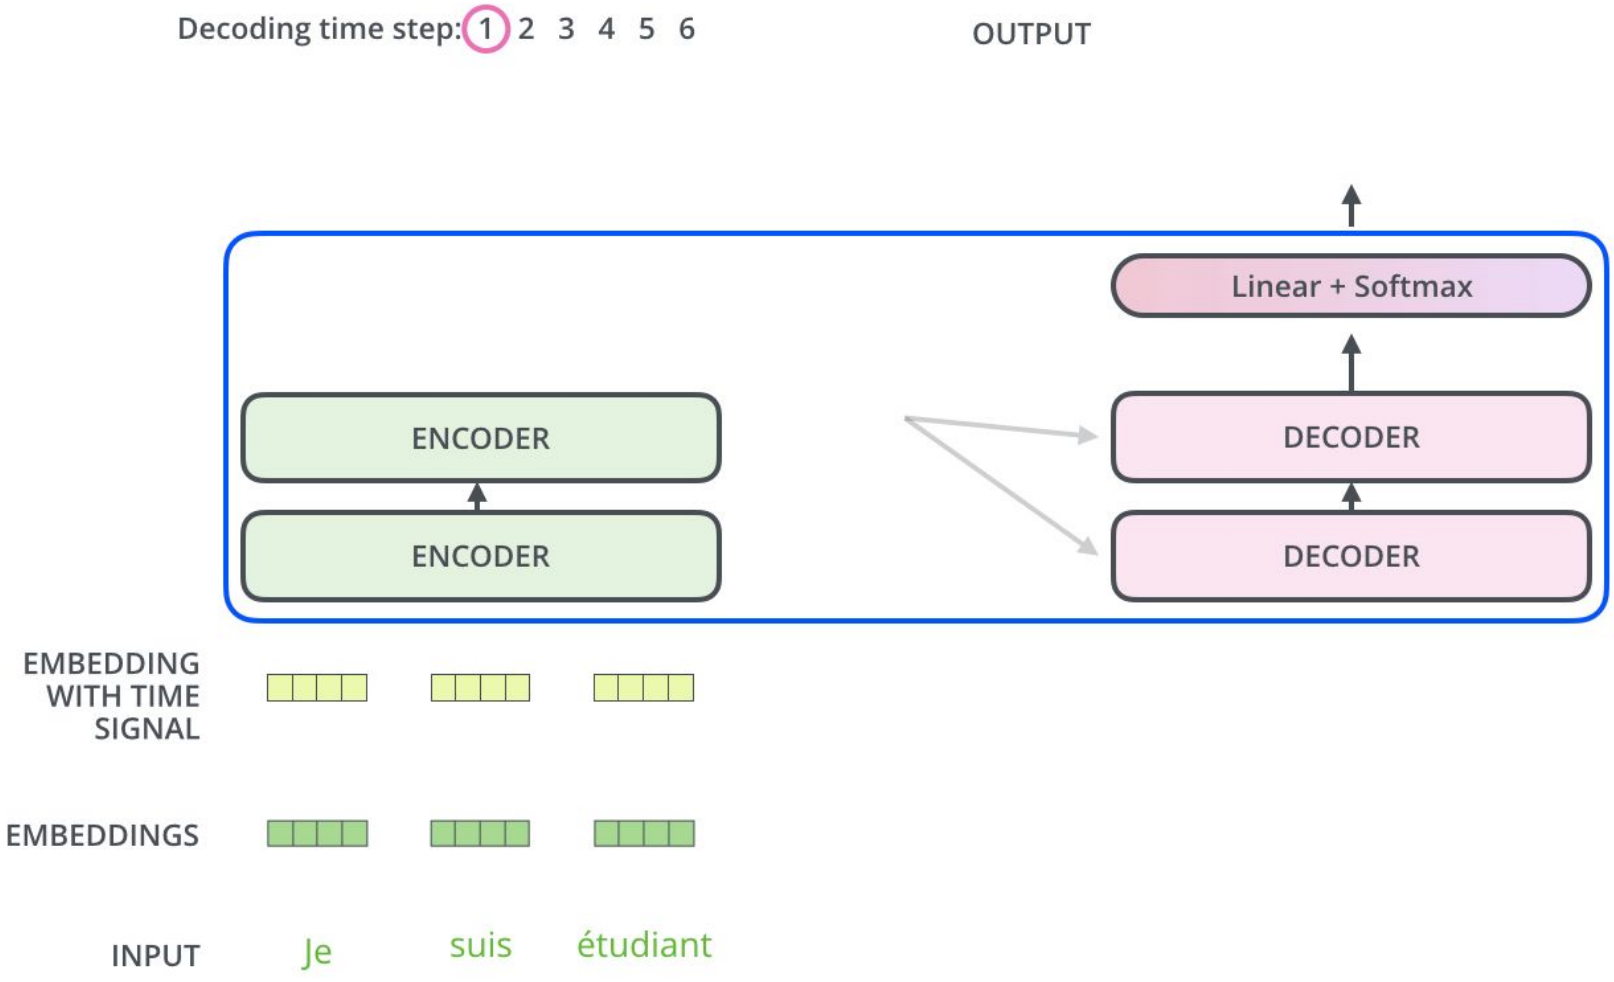

The Illustrated Transformer: http://jalammar.github.io/illustrated-transformer/

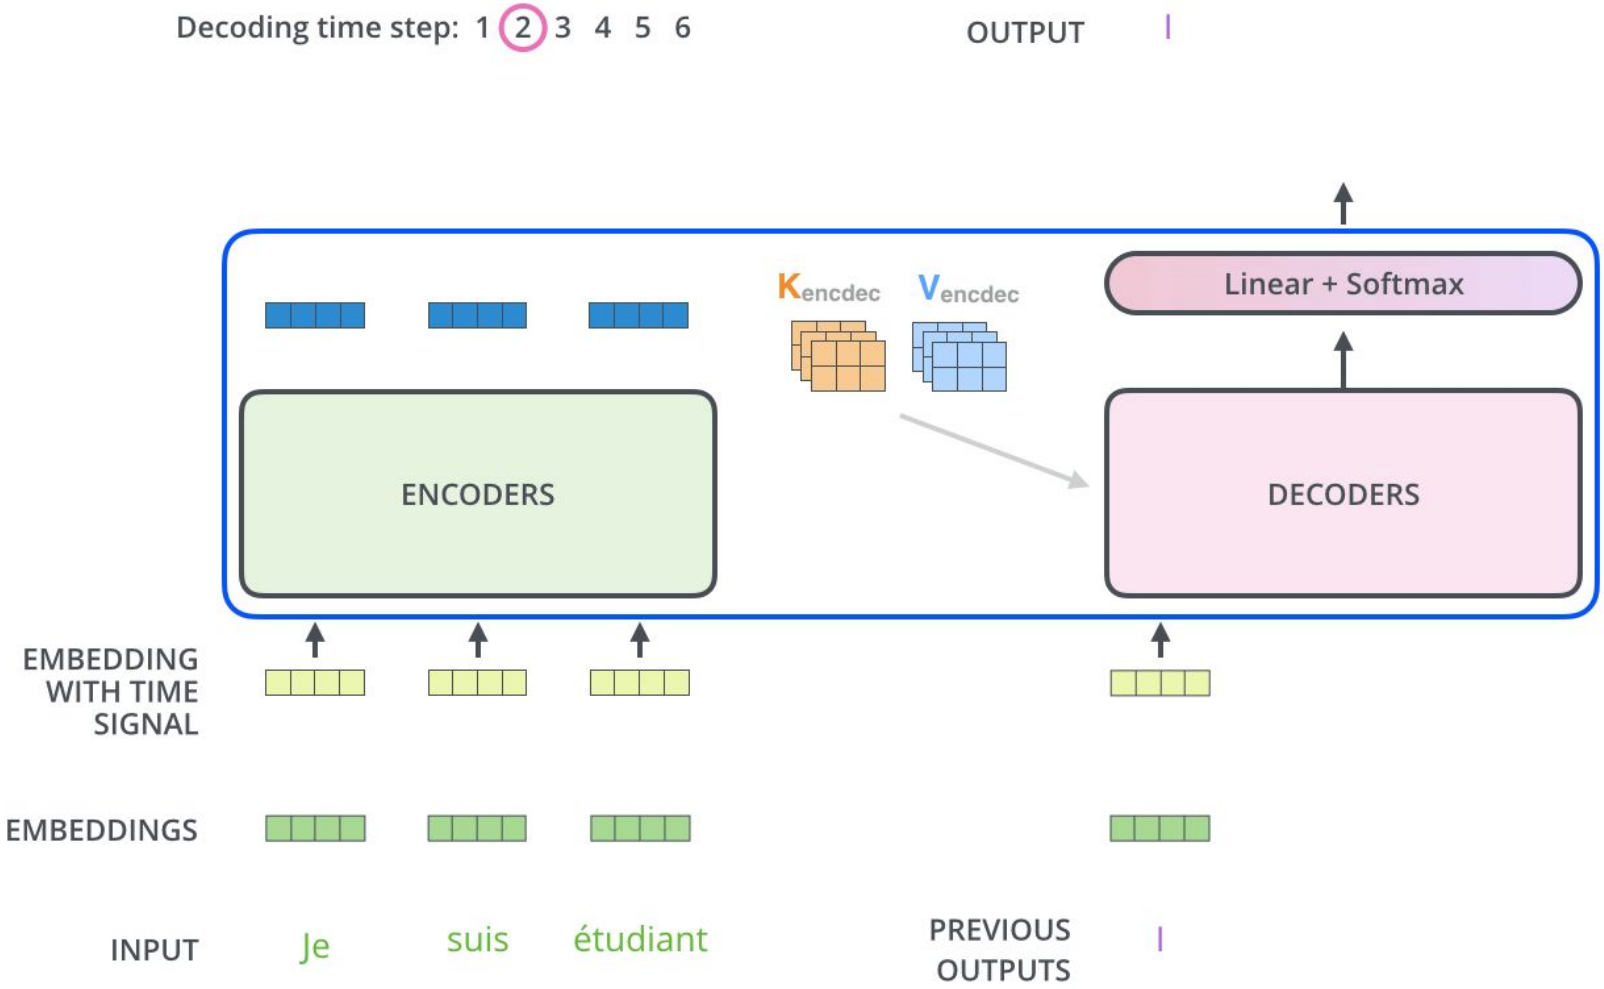

The Illustrated Transformer: http://jalammar.github.io/illustrated-transformer/

BERT Masked LM

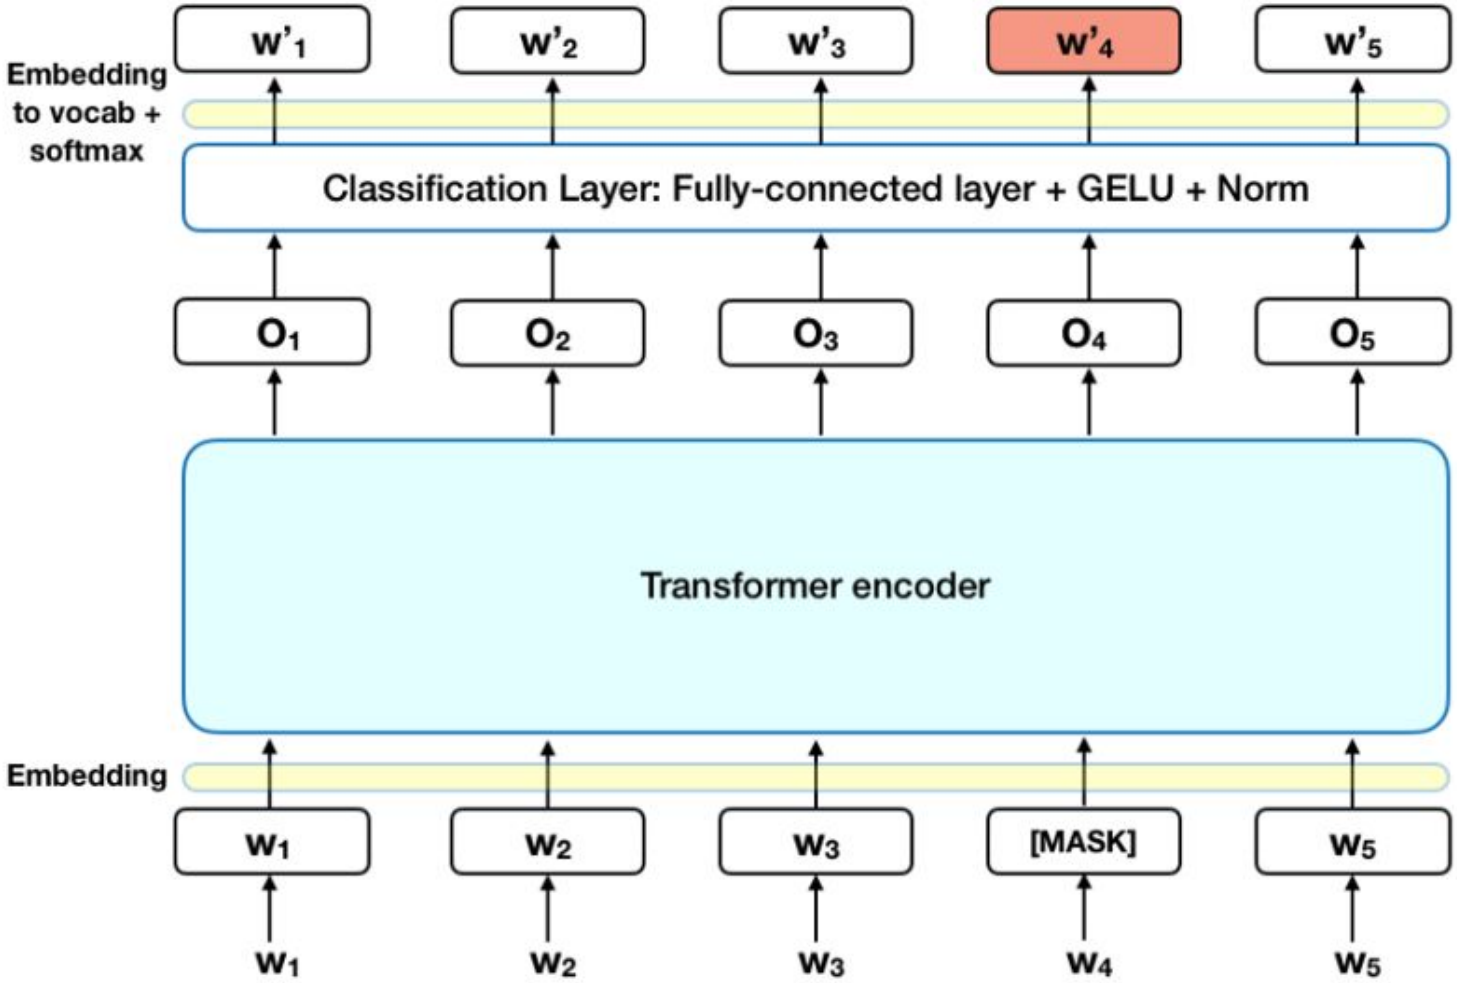

What can you do with these LMs

Part of many ML pipelines

Natural Language Understanding

Question Answering

Text Generation [https://transformer.huggingface.co/]

Topic Models

Transf

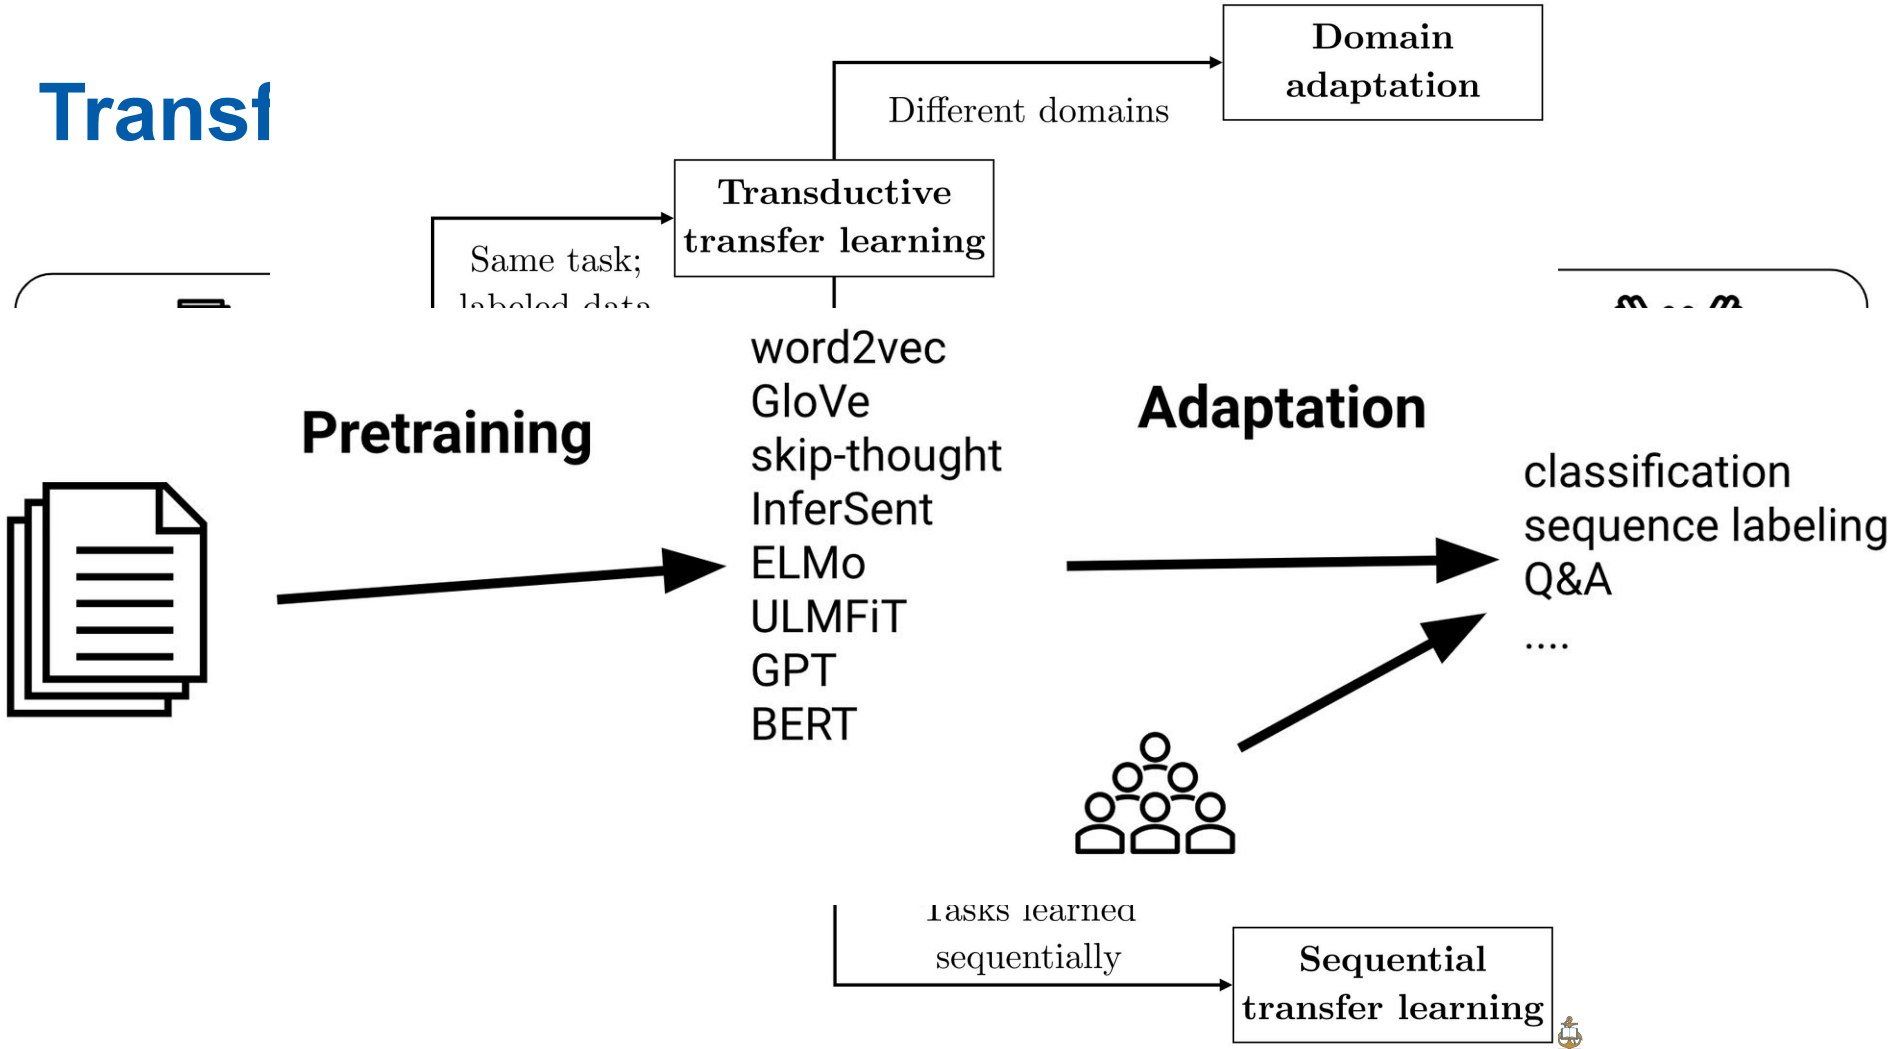

A taxonomy for transfer learning in NLP (Ruder, 2019).


GTX 1070
NVIDIA.

Resources

NLP General

- https://github.com/fastai/course-nl
- Stanford Coursera NLP Slides https://web.stanford.edu/~jurafsky/NLPCourseraSlides.html
- Sebastian Ruder Newsletter [http://ruder.io/nlp-news/]
- https://nlpprogress.com

Python Libraries

- SKLearn NLP (Working With Text Data) - URL (Nice tutorial)
- spaCY: Industrial-Strength Natural Language Processing - URL
- NLTK

Thank You

Dr. Vukosi Marivate
vukosi.marivate@cs.up.ac.za
https//dsfsi.github.io
@vukosi